



THE UNIVERSITY OF
MELBOURNE

COMP90049 Introduction to Machine Learning

K-Nearest Neighbours Classifier

Semester 2, 2020

Hadi Khorshidi, CIS

The Slides prepared by Qihong Ke, CIS


Copyright @ University of Melbourne 2020

All rights reserved. No part of the publication may be reproduced in any form by print, photoprint, microfilm or any other means without written permission from the author.

Eager Learning vs Lazy Learning




- Eager learning:
 - train a generalization model using training data
 - use the model to predict a new test sample.



Eager Learning vs Lazy Learning




- Lazy Learning (Instance-based Learning):
 - no training of the model, store training data
 - prediction: compares new testing instances with stored instances



Eager Learning vs Lazy Learning



- K-NN is a lazy learner or Instance-based Learning method.
- how to compare the training and testing data?
- how to predict the class for the testing data?





- Introduction of K-NN
- Distance Measure
 - Features
 - Distance Metrics
- K-NN for Classification
- Summary

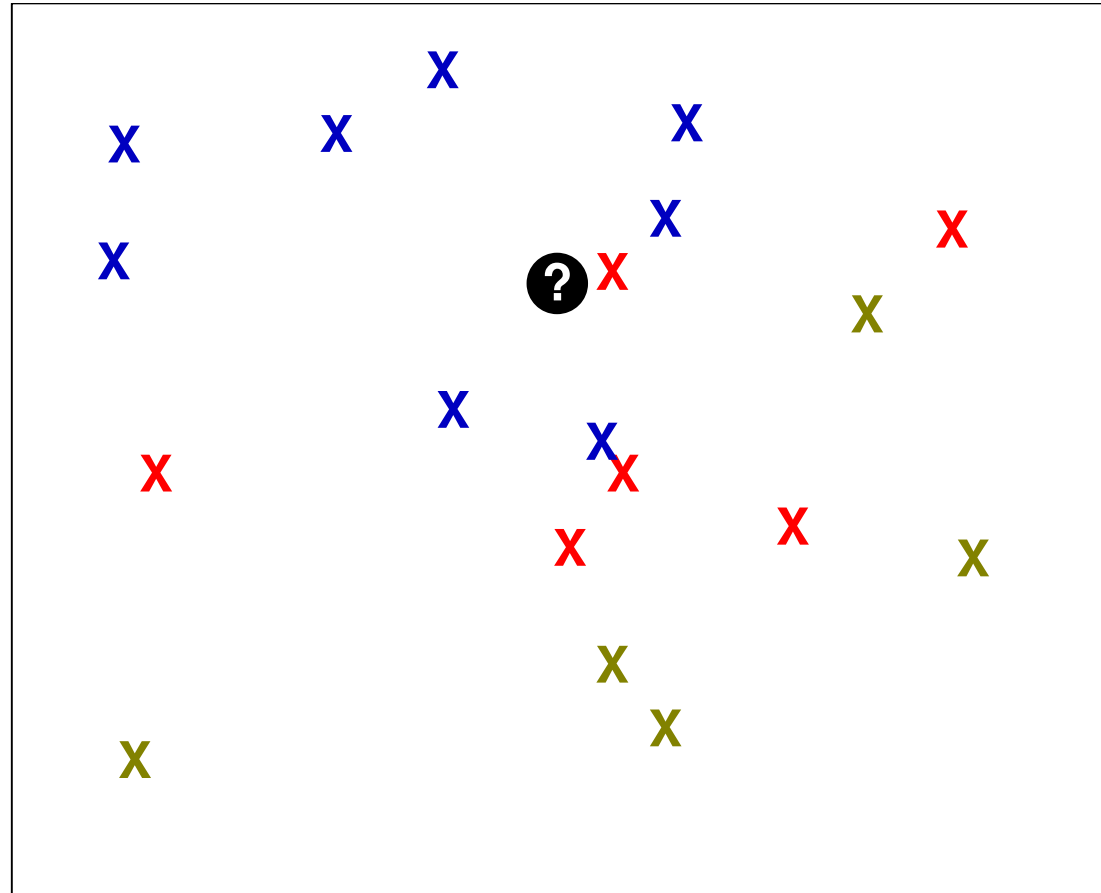


- Introduction of K-NN
- Distance Measure
 - Features
 - Distance Metrics
- K-NN for Classification
- Summary



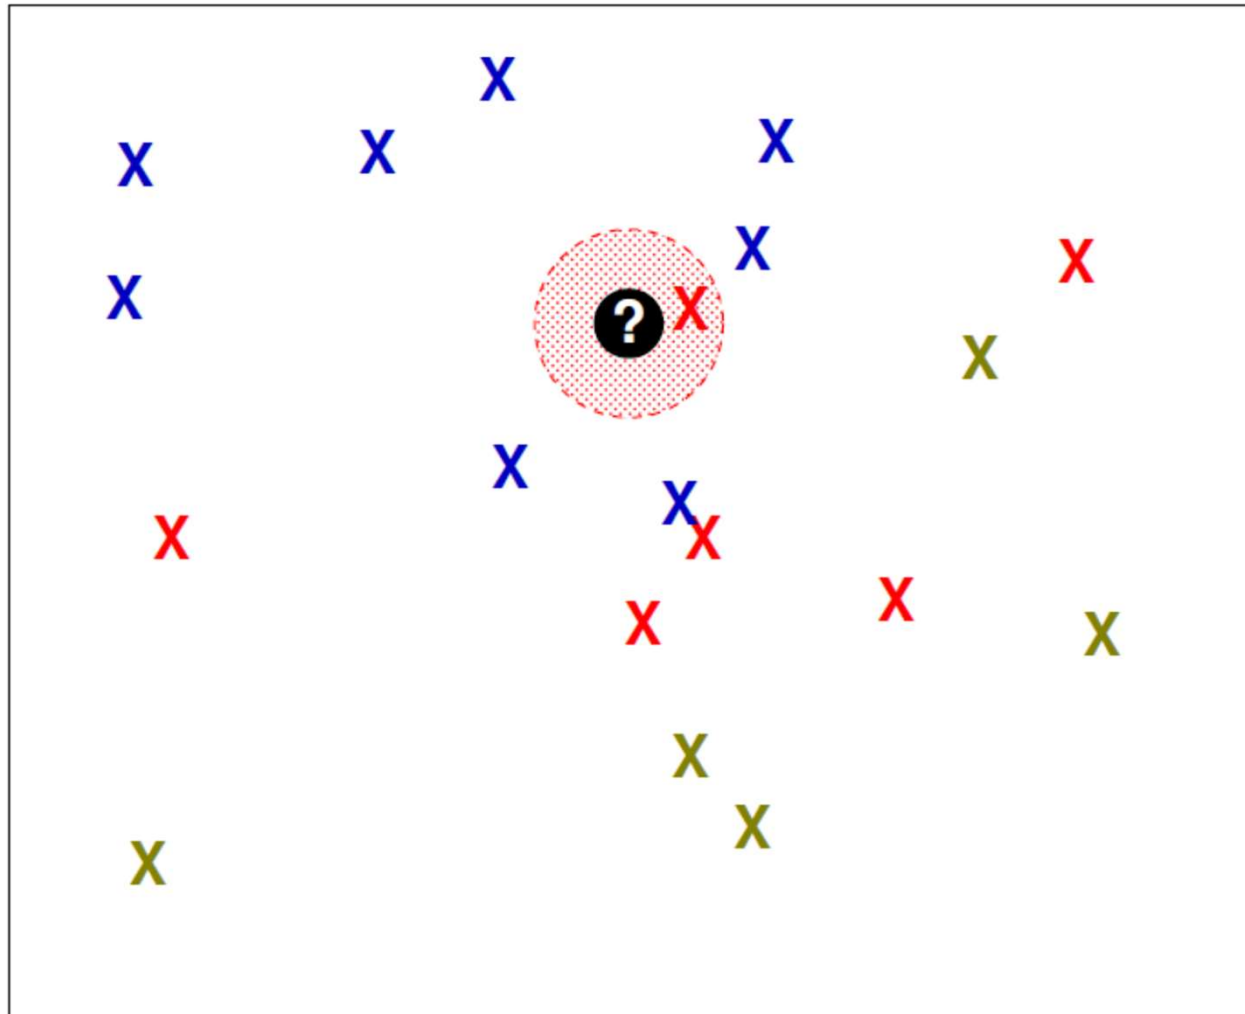
- Store training instances.
- Testing:
 - Measure the similarity (or distance) between the testing and training data
 - Find K-Nearest neighbours
 - Return the class of the testing data using the corresponding labels of the K-Nearest neighbours.

K-Nearest Neighbours

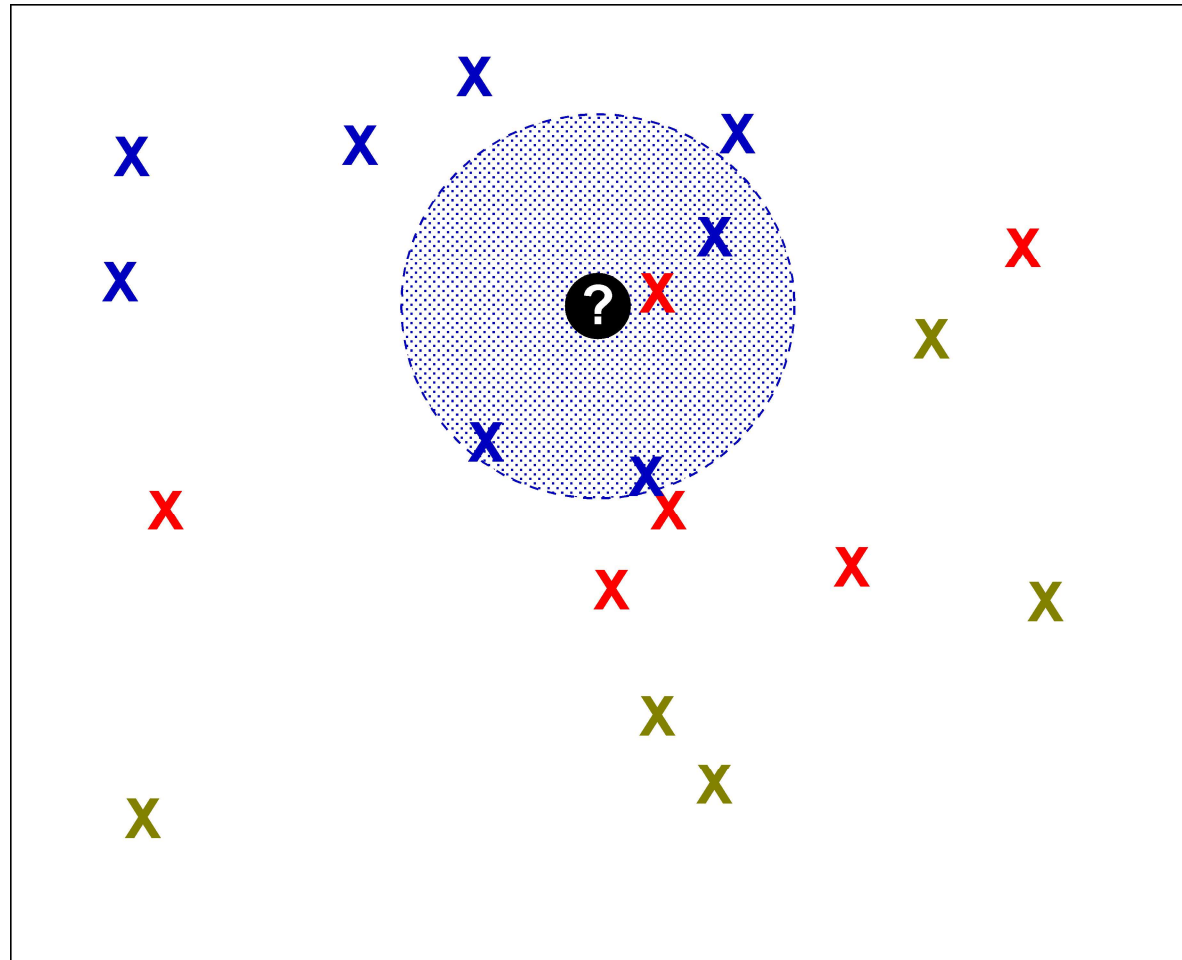


K-nearest neighbours: k closest data points

1-NN



$k=1$: the closest data point of the record



k=4: 4 closest data points of the record

Why K-NN Classifier



- Intuitive and simple
- No assumptions
- No training: new data-> evolve and adapt immediately
- used for classification and regression



- Introduction of K-NN
- Distance Measure
 - Features
 - Distance Metrics
- K-NN for Classification
- Summary

Distance vs Similarity



- Distance (Dissimilarity)
 - Numerical measure of how different two data objects are
 - Lower when objects are more alike
 - No negative distances. $d(x, y) \geq 0$
- Similarity
 - Numerical measure of how alike two data objects are.
 - Is higher when objects are more alike.

Feature Vectors



- an n-dimensional vector of features

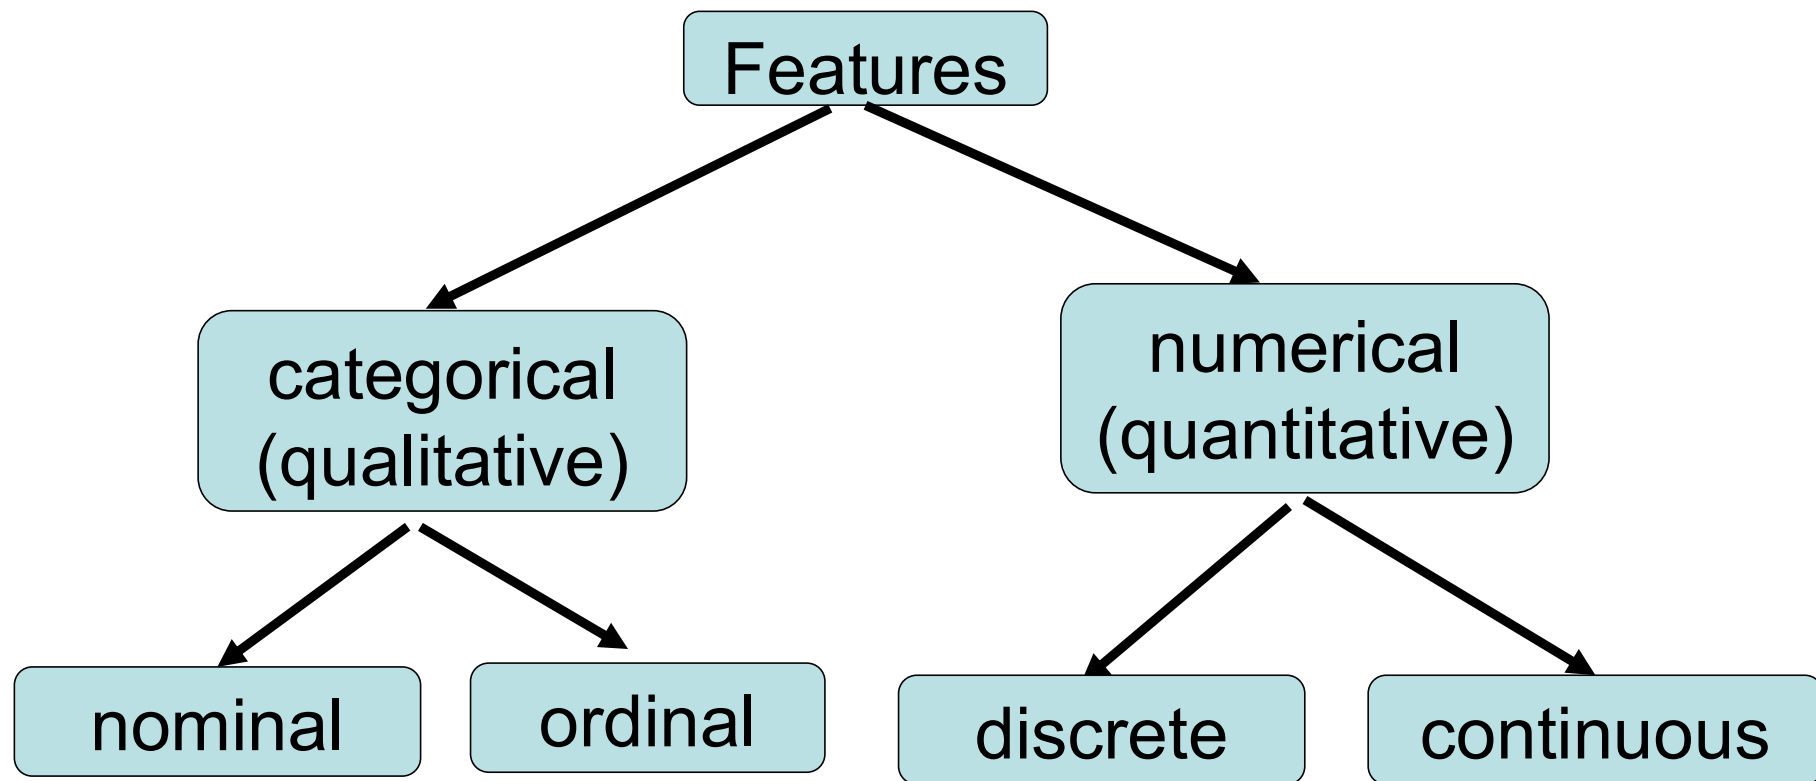
	Outlook	Temperature	Humidity	Windy	Play
a	sunny	hot	high	FALSE	no
b	sunny	hot	high	TRUE	no
c	overcast	hot	high	FALSE	yes
d	rainy	mild	high	FALSE	yes
e	rainy	cool	normal	FALSE	yes
f	rainy	cool	normal	TRUE	no

feature vector =
$$\begin{bmatrix} \text{Outlook} \\ \text{Temperature} \\ \text{Humidity} \\ \text{Windy} \end{bmatrix}$$

Features



description (characteristic, quality) of an object.





- Nominal: unordered categories, mutually exclusive,
 - colours= {black, brown, red, grey, white}
 - marital status, occupation
- Ordinal: ordered categories, mutually exclusive
 - temperature: cool < mild < hot
- Discrete: certain values, can be counted.
 - e.g., number of visits
- Continuous: any value within a range,
 - e.g., height.

Compare Nominal Feature Vectors



- Simple Matching: k : # of matches (same features), m : total # of features

$$d = \frac{m - k}{m}$$

- Example:

feature	apple	banana
shape	round	curved
colour	red	yellow
taste	sweet	sweet

$$d = \frac{3-1}{3} = \frac{2}{3}$$

Compare Nominal Feature Vectors



- Use many binary variables: one for each of the M nominal states (values)

feature	apple	banana
shape	round	curved
colour	red	yellow
taste	sweet	sweet

state	apple	banana
round	1	0
curved	0	1
red	1	0
yellow	0	1
sweet	1	1

Compare Nominal Feature Vectors



- Simple Matching: K: # of matches (same states), M: total # of nominal states

$$d = \frac{M - K}{M}$$

state	apple	banana
round	1	0
curved	0	1
red	1	0
yellow	0	1
sweet	1	1

$$dis = \frac{5 - 1}{5} = \frac{4}{5}$$

Compare Ordinal Feature Vectors

Order is important.

- sort the value, return a rank

$$r \in \{1, \dots, M\}$$

- rank -> numerical values:

$$z = \frac{r - 1}{M - 1}$$

- Compute the distance (using methods for comparing numerical variables)

feature	A	B
safety	0	2
comfortable	-2	1
convenient	-1	2

feature	A	B
safety	2/4	4/4
comfortable	0	3/4
convenient	1/4	4/4

rating:

-2: very dissatisfied. -1: dissatisfied


0: indifference. 1: satisfied.

2: very satisfied

Compare Numerical Feature Vectors : Euclidean Distance



Given two items A and B, and their feature vectors a and b , respectively, we can calculate their distance d in Euclidean space:



In n -dimensional space:

$$d(A, B) = \sqrt{\sum_{i=1}^n (a_i - b_i)^2}$$

Compare Numerical Feature Vectors : Euclidean Distance



$$d(A, B) = \sqrt{\sum_{i=1}^n (a_i - b_i)^2}$$

feature	A	B
safety	2/4	4/4
comfortable	0	3/4
convenient	1/4	4/4

$$d(A, B) = \sqrt{\left(\frac{2}{4} - 1\right)^2 + \left(0 - \frac{3}{4}\right)^2 + \left(\frac{1}{4} - 1\right)^2} = 1.17$$

Example



$$z = \frac{r - 1}{M - 1}$$

$$d(A, B) = \sqrt{\sum_{i=1}^n (a_i - b_i)^2}$$

temperature: cold<mild<hot
humidity: low<normal<high

cold, mild, hot: 1,2,3

low, normal, high: 1, 2 ,3

feature	A	B	C
temperature	hot	mild	cold
humidity	low	high	normal

feature	A	B	C
temperature	1	1/2	0
humidity	0	1	1/2


$$d(A, B) = \sqrt{(1 - \frac{1}{2})^2 + (0 - 1)^2} = 1.12$$

$$d(B, C) = \sqrt{(\frac{1}{2} - 0)^2 + (1 - \frac{1}{2})^2} = 0.71$$

Compare Numerical Feature Vectors : Manhattan Distance



["City block" distance or "Taxicab geometry" or "L1 distance"] Given two items A and B, and their corresponding feature vectors a and b , respectively, we can calculate their distance d based on their absolute differences of each feature (cartesian coordinates):



$$d(A, B) = \sum_{i=1}^n |a_i - b_i|$$

Example

$$z = \frac{r - 1}{M - 1}$$

$$d(A, B) = \sum_{i=1}^n |a_i - b_i|$$

temperature: cold < mild < hot
humidity: low < normal < high

cold, mild, hot: 1, 2, 3

low, normal, high: 1, 2, 3

feature	A	B	C
temperature	hot	mild	cold
humidity	low	high	normal

feature	A	B	C
temperature	1	1/2	0
humidity	0	1	1/2


$$d(A, B) = \left| 1 - \frac{1}{2} \right| + |0 - 1| = 1.5$$

$$d(B, C) = \left| \frac{1}{2} - 0 \right| + \left| 1 - \frac{1}{2} \right| = 1$$

Compare Numerical Feature Vectors : Cosine Similarity



- the magnitudes are ignored
- similarity range $[-1, 1]$:
 - angle 0 : similarity 1
 - angle π : similarity -1
- distance: 1-similarity: range $[0, 2]$



$$\cos(A, B) = \frac{a \cdot b}{|a||b|} = \frac{\sum_i a_i b_i}{\sqrt{\sum_i a_i^2} \sqrt{\sum_i b_i^2}}$$

Compare Numerical Feature Vectors : Cosine Similarity



used for data that the magnitude of the vectors is not important

e.g., text data represented by word counts:

- larger document-> larger word count
- topic is more important, not the size
- Using cosine similarity to ignore the magnitude.

Example



$$\cos(A, B) = \frac{a \cdot b}{|a||b|} = \frac{\sum_i a_i b_i}{\sqrt{\sum_i a_i^2} \sqrt{\sum_i b_i^2}}$$

feature	A	B	B_s
Word 1	200	300	50
Word 2	300	200	40
Word 3	200	100	25

$$\cos(A, B) = \frac{200 \times 300 + 300 \times 200 + 200 \times 100}{\sqrt{200^2 + 300^2 + 200^2} \times \sqrt{300^2 + 200^2 + 100^2}} = 0.91$$

$$\cos(B, B_s) = \frac{300 \times 50 + 200 \times 40 + 100 \times 25}{\sqrt{300^2 + 200^2 + 100^2} \times \sqrt{50^2 + 40^2 + 25^2}} = 0.99$$

Example



$$d(A, B) = \sqrt{\sum_{i=1}^n (a_i - b_i)^2}$$

$d(A, B)$


feature	A	B	B_s
Word 1	200	300	50
Word 2	300	200	40
Word 3	200	100	25

$$= \sqrt{(200 - 300)^2 + (300 - 200)^2 + (200 - 100)^2} = 173.2$$

$d(B, B_s)$

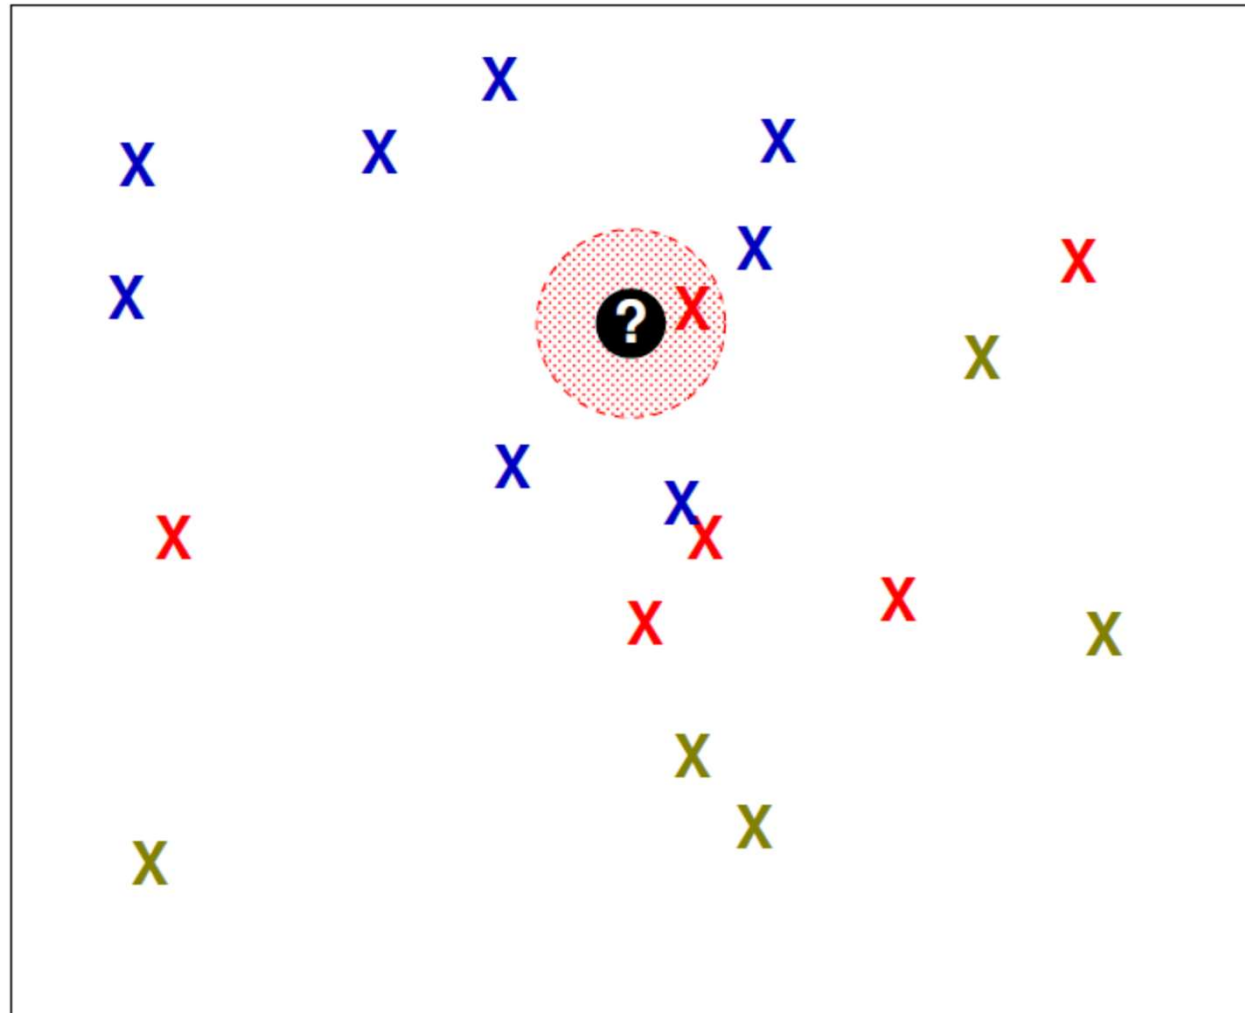
$$= \sqrt{(300 - 50)^2 + (200 - 40)^2 + (100 - 25)^2} = 306.2$$

K-NN






- Introduction of K-NN
- Distance Measure
 - Features
 - Distance Metrics
- **K-NN for Classification**
- Summary




- 1-NN: Classify the test input according to the class of the closest training instance.



- K-NN: Classify the test input according to the majority class of the K-Nearest training instances.

Weighted K-NN



- blue: $w_2 + w_3 + w_4$
- red: w_1

- weighted K-NN: Classify the test input according to the weighted accumulative class of the K-Nearest training instances



There are a number of strategies for weighting:

- Give each neighbour equal weight (= classify according to the majority class of set of neighbours)
- Weight the vote of each instance by its **inverse linear distance (ILD)** from the test instance:

$$w_j = \frac{d_k - d_j}{d_k - d_1}$$

Where

d_1 is the minimum distance between the neighbours and the test instance;

d_k is the maximum distance between the neighbours and the test instance

d_j is the distance between the j th neighbour and the test instance

- Weight the vote of each instance by its **inverse distance (ID)** from the test instance:

$$w_j = \frac{1}{d_j + \epsilon}$$

ϵ is a small constant (e.g., $1e - 10$ (10^{-10})) to make sure the denominator is not zero

- Weighted accumulative class of the K-Nearest training instances for each class:

$$c = \sum w_j$$

Voting Strategies in Action (k = 3)



Instance	Class	Distance
1	Yes	0
2	No	1
3	No	1.5

majority class voting: **no = 2** vs. yes = 1

ILD-based voting:

$$\text{yes} = \frac{1.5-0}{1.5-0} = 1$$

vs.

$$\text{no} = \frac{1.5-1}{1.5-0} + \frac{1.5-1.5}{1.5-0} = 0.33 + 0 = 0.33$$

$$w_j = \frac{d_k - d_j}{d_k - d_1}$$

Voting Strategies in Action ($k = 3$)



Instance	Class	Distance
1	Yes	0
2	No	1
3	No	1.5

ID-based voting ($\epsilon = 1e-10$):

$$\text{yes} = \frac{1}{0 + \epsilon} = 1e10$$

vs.

$$\text{no} = \frac{1}{1 + \epsilon} + \frac{1}{1.5 + \epsilon} = 1.67$$

$$w_j = \frac{1}{d_j + \epsilon}$$

Voting Strategies in Action (k = 5)



Instance	Class	Distance
1	Yes	0
2	No	1
3	No	1.5
4	Yes	1.75
5	No	2

ILD-based voting:

$$w_j = \frac{d_k - d_j}{d_k - d_1}$$

ID-based voting:

$$w_j = \frac{1}{d_j + \epsilon}$$

majority class voting: yes = 2 vs. **no = 3**

ILD-based voting:

$$\text{yes} = \frac{2-0}{2-0} + \frac{2-1.75}{2-0} = 1 + 0.125 = 1.125$$

vs.

$$\text{no} = \frac{2-1}{2-0} + \frac{2-1.5}{2-0} + \frac{2-2}{2-0} = 0.5 + 0.25 + 0 = 0.75$$

Voting Strategies in Action (k = 5)



Instance	Class	Distance
1	Yes	0
2	No	1
3	No	1.5
4	Yes	1.75
5	No	2

ILD-based voting:

$$w_j = \frac{d_k - d_j}{d_k - d_1}$$

ID-based voting:

$$w_j = \frac{1}{d_j + \epsilon}$$

ID-based voting ($\epsilon = 1e-10$):

$$\text{yes} = \frac{1}{0+\epsilon} + \frac{1}{1.75+\epsilon} = 1e10$$

vs.

$$\text{no} = \frac{1}{1+\epsilon} + \frac{1}{1.5+\epsilon} + \frac{1}{2+\epsilon} = 2.17$$



In the case that we have an equal number of votes for different classes, we need some tie breaking mechanism:

- random tie breaking
- take class with highest prior probability
- see if the addition of the $k + 1$ th instance breaks the tie

Choosing the Value of k

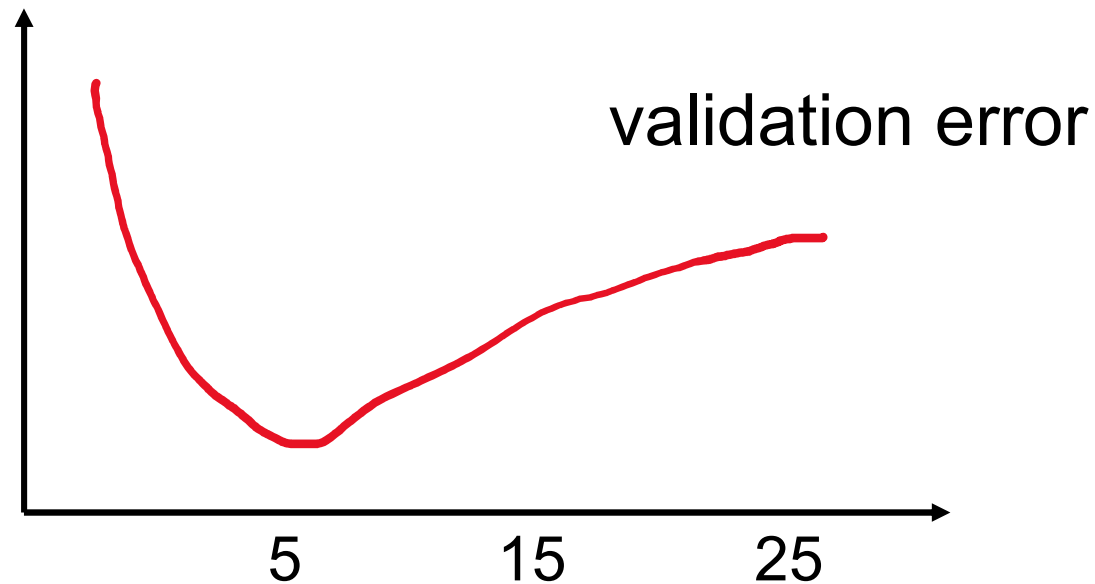


- Smaller k:
 - jagged decision boundary
 - noise
 - lower classifier performance
- Larger k:
 - smooth decision boundary,
 - includes unrelated classes
 - lower classifier performance
- $K=N$: Zero-R (simply predicts the majority class) performance

Choosing the Value of k



- Choose k by cross validation
- Note: k is generally set to an odd value



Weaknesses of NN methods

Weaknesses

- Distance function: what?
- Combining the labels of multiple neighbours: how?
- Arbitrary K value
- Expensive if data set is large.
 - Typical implementation: brute-force computation of distances between a test instance and every training instance.
 - efficient if data is small
 - infeasible if data grows



K-NN algorithm spends more computation time on testing than train.

- A) TRUE
- B) FALSE



Today: K-Nearest Neighbour

- Distance metrics for different feature types
- Euclidean distance, Manhattan distance, Cosine similarity
- Majority voting, Inverse Linear Distance (ILD) and Inverse Distance (ID) for weighted K-Nearest Neighbour classification

Next lecture:

- *Coding demo*
- *Optimization Part II*



Some slides are from:

- *Altman, Naomi S. (1992). "An introduction to kernel and nearest-neighbor nonparametric regression" (PDF). The American Statistician. 46 (3): 175-185.*
- *Data Mining: Concepts and Techniques, 2nd ed., Jiawei Han and Micheline Kamber, Morgan Kaufmann, 2006. Chapter 2.*
- *Tan et al . Introduction to Data Mining.2006. Chapter 2. Tan et al . Introduction to Data Mining. 2006. Chapter 5, Section 5.2*
- *Jeremy Nicholson & Tim Baldwin & Karin Verspoor: Machine Learning*
- <https://www.datacamp.com/community/tutorials/k-nearest-neighbor-classification-scikit-learn>
- https://en.wikipedia.org/wiki/Instance-based_learning
- <https://towardsdatascience.com/a-simple-introduction-to-k-nearest-neighbors-algorithm-b3519ed98e>
- <https://www.c-sharpcorner.com/article/knn-k-nearest-neighbors/>
- https://en.wikipedia.org/wiki/Cosine_similarity